

Springboard--Intermediate Data Science Program

Capstone Project - Data Wrangling

By Intan Sari Septiana
February, 2019

Data Source

The dataset was provided by Grab in this link: <https://www.aiforsea.com/safety>. It contains two folders, features and labels. Inside the features folder, there are 10 csv files with size 193 MB each. For the labels folder, there is one csv file with size 299 KB.

Data Wrangling Process

- First, read all csv files. For features, I concat all files into one and save it as **df**. For labels, I read it as **label** variable.
- Let's take a look on each dataframe
 - df

```
: df.head()
```

	bookingID	Accuracy	Bearing	acceleration_x	acceleration_y	acceleration_z	gyro_x	gyro_y	gyro_z	second	Speed
0	1202590843006	3.000	353.0	1.228867	8.900100	3.986968	0.008221	0.002269	-0.009966	1362.0	0.000000
1	274877907034	9.293	17.0	0.032775	8.659933	4.737300	0.024629	0.004028	-0.010858	257.0	0.190000
2	884763263056	3.000	189.0	1.139675	9.545974	1.951334	-0.006899	-0.015080	0.001122	973.0	0.667059
3	1073741824054	3.900	126.0	3.871543	10.386364	-0.136474	0.001344	-0.339601	-0.017956	902.0	7.913285
4	1056561954943	3.900	50.0	-0.112882	10.550960	-1.560110	0.130568	-0.061697	0.161530	820.0	20.419409

- label

```
label.head()
```

	bookingID	label
0	111669149733	0
1	335007449205	1
2	171798691856	0
3	1520418422900	0
4	798863917116	0

- Basically they are connected by **bookingID** columns. **df** has 16135561 rows while **label** has 20018 rows, make sense because from the data description from client, one **bookingID** has thousands telematics data.
- Check the data types of each column:
 - Label

```
: ## check data types
```

```
label.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20018 entries, 0 to 20017  
Data columns (total 2 columns):  
bookingID    20018 non-null int64  
label        20018 non-null int64  
dtypes: int64(2)  
memory usage: 312.9 KB
```

- Features

```
## check data types
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 16135561 entries, 0 to 16135560  
Data columns (total 11 columns):  
bookingID          int64  
Accuracy           float64  
Bearing            float64  
acceleration_x     float64  
acceleration_y     float64  
acceleration_z     float64  
gyro_x             float64  
gyro_y             float64  
gyro_z             float64  
second             float64  
Speed              float64  
dtypes: float64(10), int64(1)  
memory usage: 1.3 GB
```

- Looks good and **bookingID** data type matches. Let's check null values in df, since it is clear all columns in **label** contain non null value.

```
## check null value
```

```
df.isna().sum()
```

```
bookingID      0
Accuracy        0
Bearing         0
acceleration_x  0
acceleration_y  0
acceleration_z  0
gyro_x          0
gyro_y          0
gyro_z          0
second         0
Speed          0
dtype: int64
```

- No null value. Now let's check the number of unique **bookingID** in each data frame.

```
# check unique bookingID
print('features:', df.bookingID.nunique())
print('label:', label.bookingID.nunique())
```

```
features: 20000
label: 20000
```

- They have the same total of unique **bookingID**, which does not necessarily means they are fully connected (we will check this later), but what we can see here, label data frame obviously has duplicated rows since it has 20018 rows but only has 20000 **bookingID**. We know that **label** only contains a pair of **bookingID** and its label.
- Let's try to drop the duplicated rows

```
# from above we can see that label contains 18 rows with duplicate ID, let's drop the duplicates
label = label.drop_duplicates()
```

```
label.shape
```

```
(20018, 2)
```

- After trying to drop the duplicates, the rows of **label** remain the same, which means they are not duplicates. We will check what is going on here but before that we will check the relationship first.

```
label bookingID in features: bookingID      20000
dtype: int64
label bookingID not in features: bookingID      0
dtype: int64
```

- From above we can say that all **bookingID** in **label** appear in **features** which means there are no **bookingID** in **label** that is not in **features**, since 20000 is the total unique **bookingID** in **label**. This should be applied to features as well. Let's confirm it.

```
features bookingID in label: bookingID    20000
dtype: int64
features bookingID not in label: bookingID     0
dtype: int64
```

- It's confirmed. Let's get back to the duplicated **bookingID** in **label**. If we cannot drop it, it means the row is not duplicated, but the data frame only has two columns, **bookingID** and **label**. If the row is not duplicated but the **bookingID** is, the only possible reason is they have different labels, which is not good. Let's check:

	bookingID	label
12602	13	1
12463	13	0
2351	154618822837	1
5295	154618822837	0
11215	223338299461	1
6212	223338299461	0
19936	395136991308	0
6121	395136991308	1
17623	403726925929	1
8472	403726925929	0
2858	455266533495	1
10778	455266533495	0

- Since the value of **label** with the same **bookingID** is contradictory, I decided to drop all of them since we do not have the true value of each **bookingID**. Now **label** only contains unique rows with unique **bookingID**

```
label.shape
```

```
(19982, 2)
```

```
print('label: ', label.bookingID.nunique())
```

```
label: 19982
```

- These changes should affect features too, so I decided to merge them to one data frame so deleted **bookingID** will be excluded on features and also for easier processing later.

```
print('shape: ',df.shape)
print('unique id: ',df.bookingID.nunique())
```

```
shape: (16116704, 12)
unique id: 19982
```

- This clean dataset can be used for later processing